Agent Evals Platform — Take-Home Assignment

Stack: NestJS (backend) · Vue 3 + TypeScript (frontend)

Problem
We build AI agents. When we change the model, add a skill, or turn on memory, we need to know whether the agent got better, got worse, or stayed the same.

Build a platform that answers this with evidence. Build it on top of Hermes Agent (https://github.com/NousResearch/hermes-agent), an open-source agent harness with skills, memory and tools.

What it should do

- Define a set of tasks, and how each task is scored.
- Define agent configurations to compare, such as model, skills, memory and tools.
- Run the tasks under each configuration.
- Show whether one configuration is better than another, and how confident we can be in that result.
- Let a user inspect an individual run to understand why it passed or failed.

Requirements

- NestJS backend, Vue 3 frontend, TypeScript throughout.
- Runs locally with a single command
- Includes a sample task set that demonstrates the platform.
- Document any assumptions you make.

Deliverables

1. Production-grade source code in a Git repository, with a README.
2. A design document covering the architecture and key decisions.
3. The known limits of the system.
4. Approximate costs of running the system.
5. Results of one real experiment comparing two configurations.
6. A video of the working system.